

**PARALLEL AND DISTRIBUTIVE COMPUTING (UCS645)**

**Lab – 1**

**By**

**Mannya Soni**

**102316067**

**3P13**



**THAPAR INSTITUTE  
OF ENGINEERING & TECHNOLOGY  
(Deemed to be University)**

**Instructor: Dr. Saif Nalband**

**THAPAR INSTITUTE OF ENGINEERING AND  
TECHNOLOGY, (A DEEMED TO BE UNIVERSITY),  
PATIALA, PUNJAB  
INDIA**

**Session-Year (Jan-May, 2026)**

## Q1. DAXPY Loop Analysis:

```
GNU nano 7.2                                     q1_daxpy_serial.c *
```

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>

int main() {
    const int N = 1 << 16;
    const double a = 2.5;
    double *X = (double*)malloc(N * sizeof(double));
    double *Y = (double*)malloc(N * sizeof(double));
    if (!X || !Y) { printf("Malloc failed\n"); return 1; }

    for (int i = 0; i < N; i++) {
        X[i] = 1.0 + i * 0.001;
        Y[i] = 2.0 - i * 0.0005;
    }

    double start = omp_get_wtime();

    for (int i = 0; i < N; i++) {
        X[i] = a * X[i] + Y[i];
    }

    double end = omp_get_wtime();
    printf("Serial time: %f seconds\n", end - start);

    free(X);
    free(Y);
    return 0;
}
```

^G Help ^O Write Out ^W Where Is ^K Cut [ Cancelled ] ^C Location M-U Undo M-A Set Mark M-J To Bracket  
^X Exit ^R Read File ^N Replace ^U Paste ^T Execute ^J Justify ^Y Go To Line M-E Redo M-G Copy M-Q Where Was

```
root@Manny-Soni:~/lab1  x  +  v
root@Manny-Soni:~# cd lab1
root@Manny-Soni:~/lab1# ls
mm_s_pi_s ql_daxpy_serial.c qls q2_mm_serial.c q3_pi_serial.c
root@Manny-Soni:~/lab1# nano ql_daxpy_serial.c
root@Manny-Soni:~/lab1# ./qls
Serial time: 0.000037 seconds
root@Manny-Soni:~/lab1#
```

Code Blame 22 lines (22 loc) · 321 Bytes

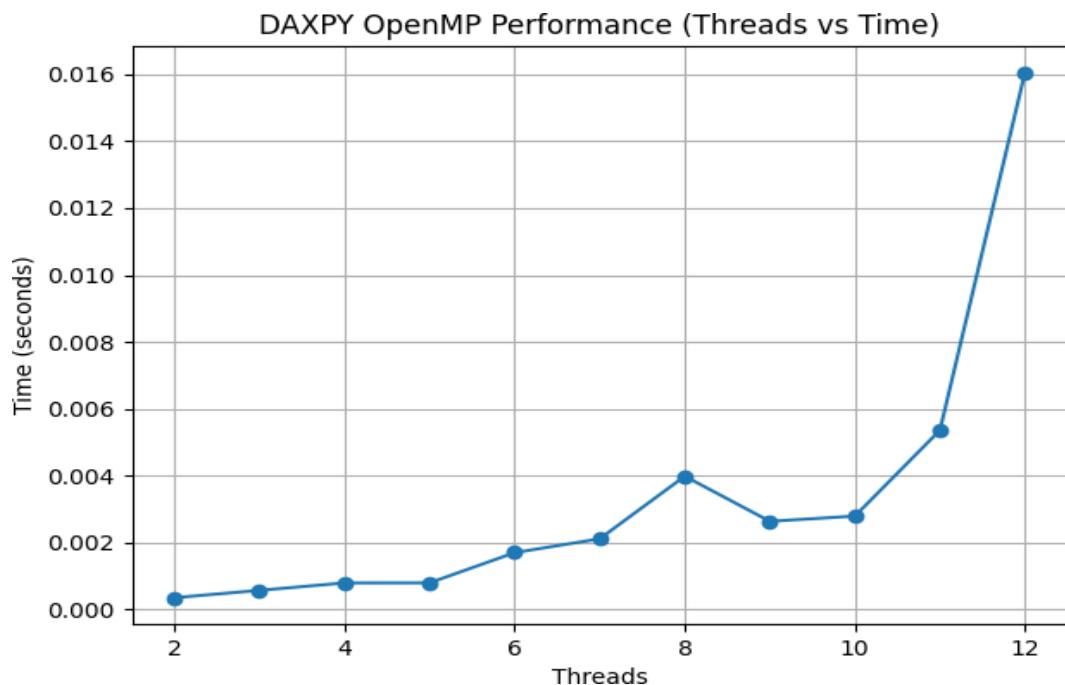
```
1  #include <stdio.h>
2  #include <omp.h>
3  #define N (1<<16)
4  int main(){
5      double x[N],y[N];
6      double a=5;
7      int i;
8      for(i=0;i<N;i++)
9      {
10         x[i]=i;
11         y[i]=i+1;
12     }
13     double start=omp_get_wtime();
14     #pragma omp parallel for
15     for(i=0;i<N;i++)
16     {
17         x[i] = a * x[i] + y[i];
18     }
19     double end=omp_get_wtime();
20     printf("time taken=%f", end - start);
21     return 0;
22 }
```

Raw ⌂ ⌄ ⌅ ⌆ ⌇ ⌈ ⌉ ⌊ ⌋ ⌃ ⌁ ⌂ ⌄ ⌅ ⌆ ⌇ ⌈ ⌉ ⌊ ⌋ ⌃ ⌁

```

root@Mannya-Soni:~/Lab1# nano q1_daxpy_omp.c
root@Mannya-Soni:~/Lab1# gcc q1_daxpy_omp.c -fopenmp -O2 -o q1p
root@Mannya-Soni:~/Lab1# OMP_NUM_THREADS=2 ./q1p
time taken=0.000341root@Mannya-Soni:~/Lab1# OMP_NUM_THREADS=3 ./q1p
time taken=0.000560root@Mannya-Soni:~/Lab1# OMP_NUM_THREADS=4 ./q1p
time taken=0.000783root@Mannya-Soni:~/Lab1# OMP_NUM_THREADS=5 ./q1p
time taken=0.000785root@Mannya-Soni:~/Lab1# OMP_NUM_THREADS=6 ./q1p
time taken=0.001689root@Mannya-Soni:~/Lab1# OMP_NUM_THREADS=7 ./q1p
time taken=0.002103root@Mannya-Soni:~/Lab1# OMP_NUM_THREADS=8 ./q1p
time taken=0.003977root@Mannya-Soni:~/Lab1# OMP_NUM_THREADS=9 ./q1p
time taken=0.002626root@Mannya-Soni:~/Lab1#
root@Mannya-Soni:~/Lab1# OMP_NUM_THREADS=10 ./q1p
time taken=0.002783root@Mannya-Soni:~/Lab1# OMP_NUM_THREADS=11 ./q1p
time taken=0.005353root@Mannya-Soni:~/Lab1# OMP_NUM_THREADS=12 ./q1p
time taken=0.016066root@Mannya-Soni:~/Lab1#

```



### Observation :

From the execution time graph, it is observed that increasing the number of threads does not significantly improve performance. In fact, after a small number of threads, the execution time increases. This happens because DAXPY is a memory-bound operation where threads compete for memory bandwidth and cache resources. Therefore, DAXPY shows poor scalability with OpenMP and performs best with fewer threads.

## **Q2. Matrix Multiplication Analysis:**

```
root@Manny-Soni:~/lab1# nano q2_mm_serial.c
GNU nano 7.2
q2_mm_serial.c

#include <stdio.h>
#include <stdlib.h>
#include <omp.h>

int main() {
    const int N = 1000;
    double *A = (double*)malloc(N*N*sizeof(double));
    double *B = (double*)malloc(N*N*sizeof(double));
    double *C = (double*)malloc(N*N*sizeof(double));
    if (!A || !B || !C) { printf("Malloc failed\n"); return 1; }

    for (int i = 0; i < N*N; i++) {
        A[i] = 1.0;
        B[i] = 1.0;
        C[i] = 0.0;
    }

    double start = omp_get_wtime();

    for (int i = 0; i < N; i++) {
        for (int j = 0; j < N; j++) {
            double sum = 0.0;
            for (int k = 0; k < N; k++) {
                sum += A[i*N + k] * B[k*N + j];
            }
            C[i*N + j] = sum;
        }
    }

    double end = omp_get_wtime();
    printf("Serial MM time: %f sec\n", end - start);

    free(A); free(B); free(C);
    return 0;
}

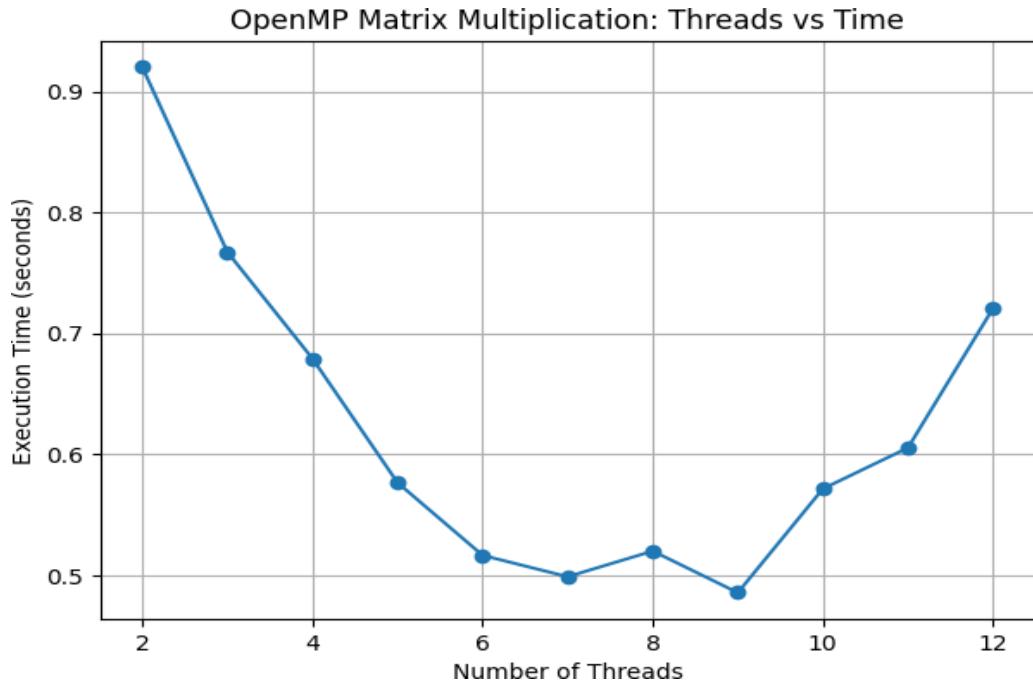
[G] Help          [^C] Write Out      [^W] Where Is      [^K] Cut           [^T] Execute      [^C] Location      [M-U] Undo
[X] Exit          [^R] Read File       [^V] Replace        [^U] Paste         [^J] Justify       [^G] Go To Line    [M-E] Redo
[M-A] Set Mark    [M-J] To Bracket   [M-C] Copy          [^Q] Where Was
```

Code Blame 33 lines (31 loc) · 756 Bytes

Raw ⌂ ⌄ ⌅ ⌆

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <omp.h>
4 #define N 1000
5
6 int main() {
7     int i, j, k;
8     double *A, *B, *C;
9     A = (double*)malloc(N * N * sizeof(double));
10    B = (double*)malloc(N * N * sizeof(double));
11    C = (double*)malloc(N * N * sizeof(double));
12    for (i = 0; i < N*N; i++) {
13        A[i] = 2.0;
14        B[i] = 3.0;
15        C[i] = 0.0;
16    }
17    double start = omp_get_wtime();
18    #pragma omp parallel for
19    for (i = 0; i < N; i++) {
20        for (j = 0; j < N; j++) {
21            for (k = 0; k < N; k++) {
22                C[i*N + j] += A[i*N + k] * B[k*N + j];
23            }
24        }
25    }
26    double end = omp_get_wtime();
27    printf("Time taken = %f", end - start);
28    free(A);
29    free(B);
30    free(C);
31    return 0;
32 }
```

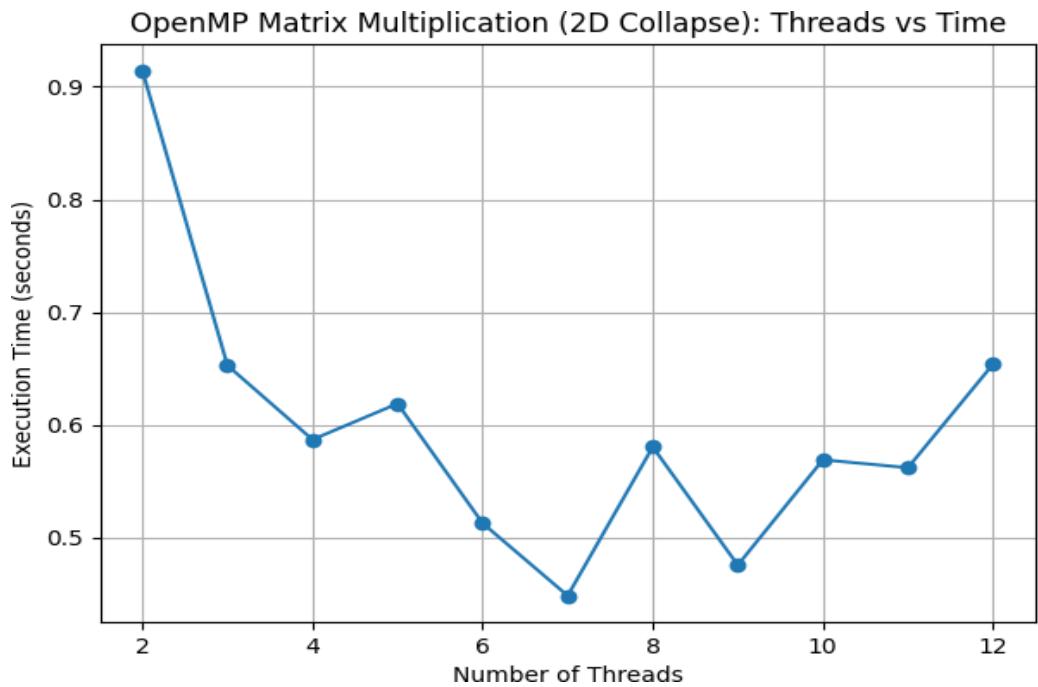
```
root@Manya-Soni:~/lab1# nano q2_mm_1d.c
root@Manya-Soni:~/lab1# gcc q2_mm_1d.c -fopenmp -O2 -o mm_1d
root@Manya-Soni:~/lab1# OMP_NUM_THREADS=2 ./mm_1d
Time taken = 0.919969root@Manya-Soni:~/Lab1# OMP_NUM_THREADS=3 ./mm_1d
Time taken = 0.767318root@Manya-Soni:~/Lab1# OMP_NUM_THREADS=4 ./mm_1d
Time taken = 0.679285root@Manya-Soni:~/Lab1# OMP_NUM_THREADS=5 ./mm_1d
Time taken = 0.576821root@Manya-Soni:~/Lab1# ^C
root@Manya-Soni:~/lab1# OMP_NUM_THREADS=6 ./mm_1d
Time taken = 0.516749root@Manya-Soni:~/Lab1# OMP_NUM_THREADS=7 ./mm_1d
Time taken = 0.499091root@Manya-Soni:~/Lab1# OMP_NUM_THREADS=8 ./mm_1d
Time taken = 0.502290root@Manya-Soni:~/Lab1# OMP_NUM_THREADS=9 ./mm_1d
Time taken = 0.485900root@Manya-Soni:~/Lab1# OMP_NUM_THREADS=10 ./mm_1d
Time taken = 0.571730root@Manya-Soni:~/Lab1# OMP_NUM_THREADS=11 ./mm_1d
Time taken = 0.605598root@Manya-Soni:~/Lab1# OMP_NUM_THREADS=12 ./mm_1d
Time taken = 0.720786root@Manya-Soni:~/lab1#
```



```

root@Mannya-Soni:~/lab1# touch q2_mm_2d.c
root@Mannya-Soni:~/lab1# touch q2_mm_2d.c
root@Mannya-Soni:~/lab1# nano q2_mm_2d.c
root@Mannya-Soni:~/lab1# gcc q2_mm_2d.c -fopenmp -O2 -o mm_2d
root@Mannya-Soni:~/lab1# OMP_NUM_THREADS=2 ./mm_2d
2D Threading Time = 0.914642root@Mannya-Soni:~/lab1# OMP_NUM_THREADS=3 ./mm_2d
2D Threading Time = 0.652942root@Mannya-Soni:~/lab1# OMP_NUM_THREADS=4 ./mm_2d
2D Threading Time = 0.587082root@Mannya-Soni:~/lab1# OMP_NUM_THREADS=5 ./mm_2d
2D Threading Time = 0.618995root@Mannya-Soni:~/lab1# OMP_NUM_THREADS=6 ./mm_2d
2D Threading Time = 0.512936root@Mannya-Soni:~/lab1# OMP_NUM_THREADS=7 ./mm_2d
2D Threading Time = 0.448649root@Mannya-Soni:~/lab1# OMP_NUM_THREADS=8 ./mm_2d
2D Threading Time = 0.580598root@Mannya-Soni:~/lab1# OMP_NUM_THREADS=9 ./mm_2d
2D Threading Time = 0.475766root@Mannya-Soni:~/lab1# OMP_NUM_THREADS=10 ./mm_2d
2D Threading Time = 0.569172root@Mannya-Soni:~/lab1# OMP_NUM_THREADS=11 ./mm_2d
2D Threading Time = 0.562154root@Mannya-Soni:~/lab1# OMP_NUM_THREADS=12 ./mm_2d
2D Threading Time = 0.653956root@Mannya-Soni:~/lab1#

```



## Observation:

The results show a clear reduction in execution time as the number of threads increases, demonstrating good parallel scalability. Since matrix multiplication is compute-intensive, the workload is effectively distributed across threads, leading to significant speedup. The 2D collapse strategy performs better than 1D threading because it improves load balancing and reduces idle time among threads. However, after a certain thread, performance begins to degrade due to overheads.

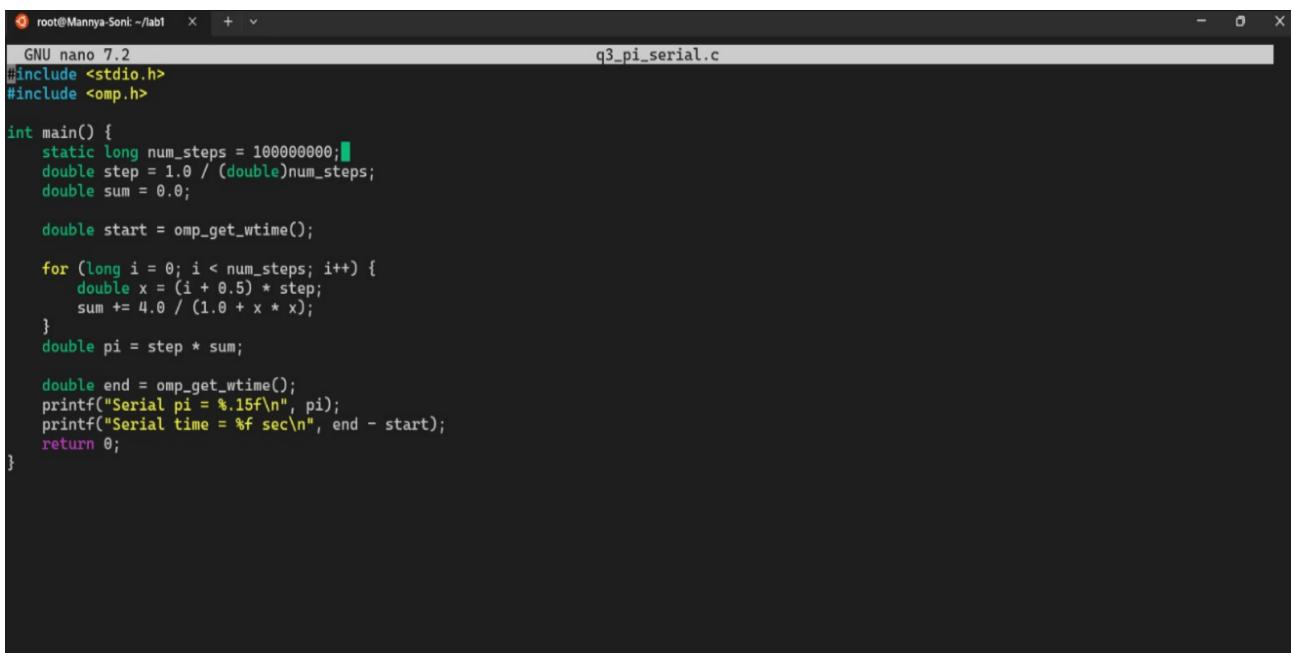
### Q3. $\pi$ Calculation Analysis:

#### Problem Statement:

The value of  $\pi$  is computed using numerical integration:

$$\pi = \int_0^1 \frac{4}{1+x^2} dx$$

The integral is approximated using the rectangle method and parallelized using OpenMP reduction.



```
root@Mannya-Soni:~/lab1  × + ▾          q3_pi_serial.c
GNU nano 7.2
#include <stdio.h>
#include <omp.h>

int main() {
    static long num_steps = 100000000;
    double step = 1.0 / (double)num_steps;
    double sum = 0.0;

    double start = omp_get_wtime();

    for (long i = 0; i < num_steps; i++) {
        double x = (i + 0.5) * step;
        sum += 4.0 / (1.0 + x * x);
    }
    double pi = step * sum;

    double end = omp_get_wtime();
    printf("Serial pi = %.15f\n", pi);
    printf("Serial time = %f sec\n", end - start);
    return 0;
}
```



```
root@Mannya-Soni:~/lab1  × + ▾
root@Mannya-Soni:~/lab1# nano q3_pi_serial.c
root@Mannya-Soni:~/lab1# ./mm_s
Serial MM time: 0.906190 sec
root@Mannya-Soni:~/lab1#
```

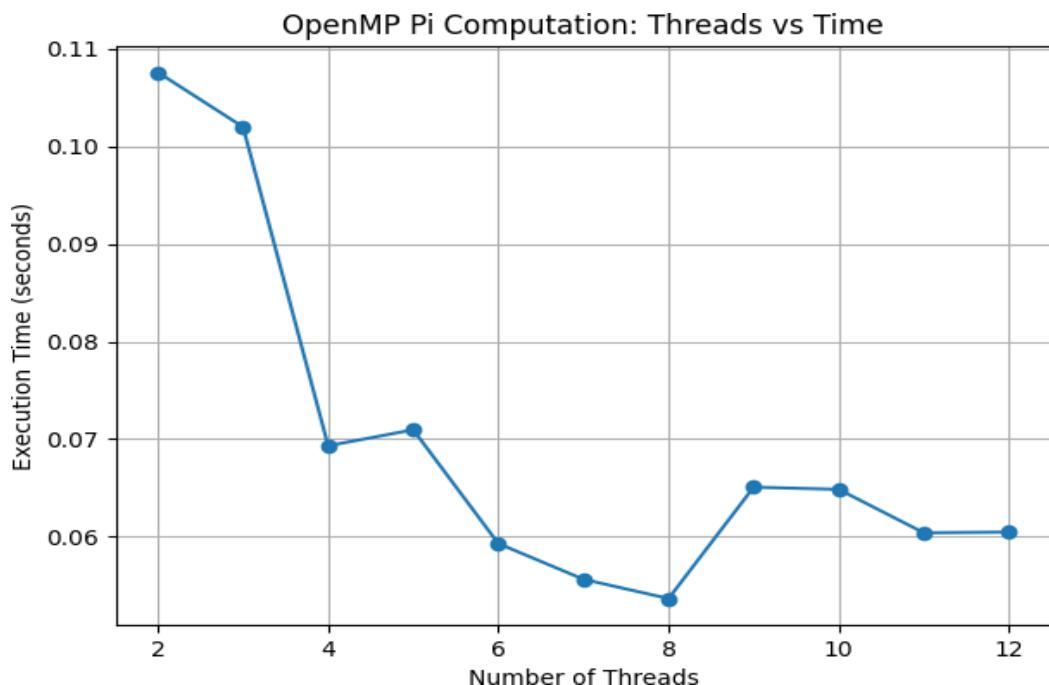
Code Blame 19 lines (19 loc) · 523 Bytes

```

1  #include <stdio.h>
2  #include <omp.h>
3  #define STEPS 100000000
4  int main() {
5      int i;
6      double x, pi, sum = 0.0;
7      double step = 1.0 / STEPS;
8      double start = omp_get_wtime();
9      #pragma omp parallel for private(x) reduction(+:sum)
10     for (i = 0; i < STEPS; i++) {
11         x = (i + 0.5) * step;
12         sum += 4.0 / (1.0 + x * x);
13     }
14     pi = step * sum;
15     double end = omp_get_wtime();
16     printf("Calculated value of Pi = %.10f\n", pi);
17     printf("Time taken = %f seconds\n", end - start);
18     return 0;
19 }
```

```

root@Mannya-Soni:~/lab1# nano q3_pi_omp.c
root@Mannya-Soni:~/lab1# gcc q3_pi_omp.c -fopenmp -O2 -o pi_p
root@Mannya-Soni:~/lab1# OMP_NUM_THREADS=2 ./pi_p
Calculated value of Pi = 3.1415926536
Time taken = 0.187599 seconds
root@Mannya-Soni:~/lab1# OMP_NUM_THREADS=3 ./pi_p
Calculated value of Pi = 3.1415926536
Time taken = 0.102003 seconds
root@Mannya-Soni:~/lab1# OMP_NUM_THREADS=4 ./pi_p
Calculated value of Pi = 3.1415926536
Time taken = 0.069321 seconds
root@Mannya-Soni:~/lab1# OMP_NUM_THREADS=5 ./pi_p
Calculated value of Pi = 3.1415926536
Time taken = 0.070958 seconds
root@Mannya-Soni:~/lab1# OMP_NUM_THREADS=6 ./pi_p
Calculated value of Pi = 3.1415926536
Time taken = 0.059295 seconds
root@Mannya-Soni:~/lab1# OMP_NUM_THREADS=7 ./pi_p
Calculated value of Pi = 3.1415926536
Time taken = 0.055633 seconds
root@Mannya-Soni:~/lab1# OMP_NUM_THREADS=8 ./pi_p
Calculated value of Pi = 3.1415926536
Time taken = 0.053658 seconds
root@Mannya-Soni:~/lab1# OMP_NUM_THREADS=9 ./pi_p
Calculated value of Pi = 3.1415926536
Time taken = 0.065079 seconds
root@Mannya-Soni:~/lab1# OMP_NUM_THREADS=10 ./pi_p
Calculated value of Pi = 3.1415926536
Time taken = 0.064851 seconds
root@Mannya-Soni:~/lab1# OMP_NUM_THREADS=11 ./pi_p
Calculated value of Pi = 3.1415926536
Time taken = 0.066395 seconds
root@Mannya-Soni:~/lab1# OMP_NUM_THREADS=12 ./pi_p
Calculated value of Pi = 3.1415926536
Time taken = 0.066477 seconds
root@Mannya-Soni:~/lab1#
```



**Observation:**

The parallel reduction approach significantly improves performance compared to the serial implementation. Execution time decreases steadily as the number of threads increases up to the number of available CPU cores. Beyond this point, the performance gain saturates and may slightly degrade due to synchronization overhead and thread management costs.